

# UNCERTAINTY-AWARE IMITATION LEARNING FOR SAFE INDOOR ROBOT NAVIGATION

Shubham Yogesh Jangle

Master's Student, Robotics Engineering University of California, Riverside

Advisor: Jiachen Li Committee: Ioannis Karamouzas

## THE CHALLENGE: WHEN ROBOTS DON'T KNOW WHAT THEY DON'T KNOW

- **Traditional Imitation Learning Problem:**

Robots learn from demonstrations in one environment

Deploy in slightly different environment

Policy produces confident predictions... even when wrong

Result: Collisions OR overly conservative behavior

- **Key Question:**

"How can robots recognize when their predictions are unreliable?"

# PROJECT GOALS

## Main Goal:

Enable safe autonomous navigation with uncertainty awareness

## Four Specific Objectives:

I. Detect when policy predictions are unreliable

→ Monte Carlo Dropout uncertainty estimation

3. Improve performance through targeted learning

→ DAgger with uncertainty-guided data collection

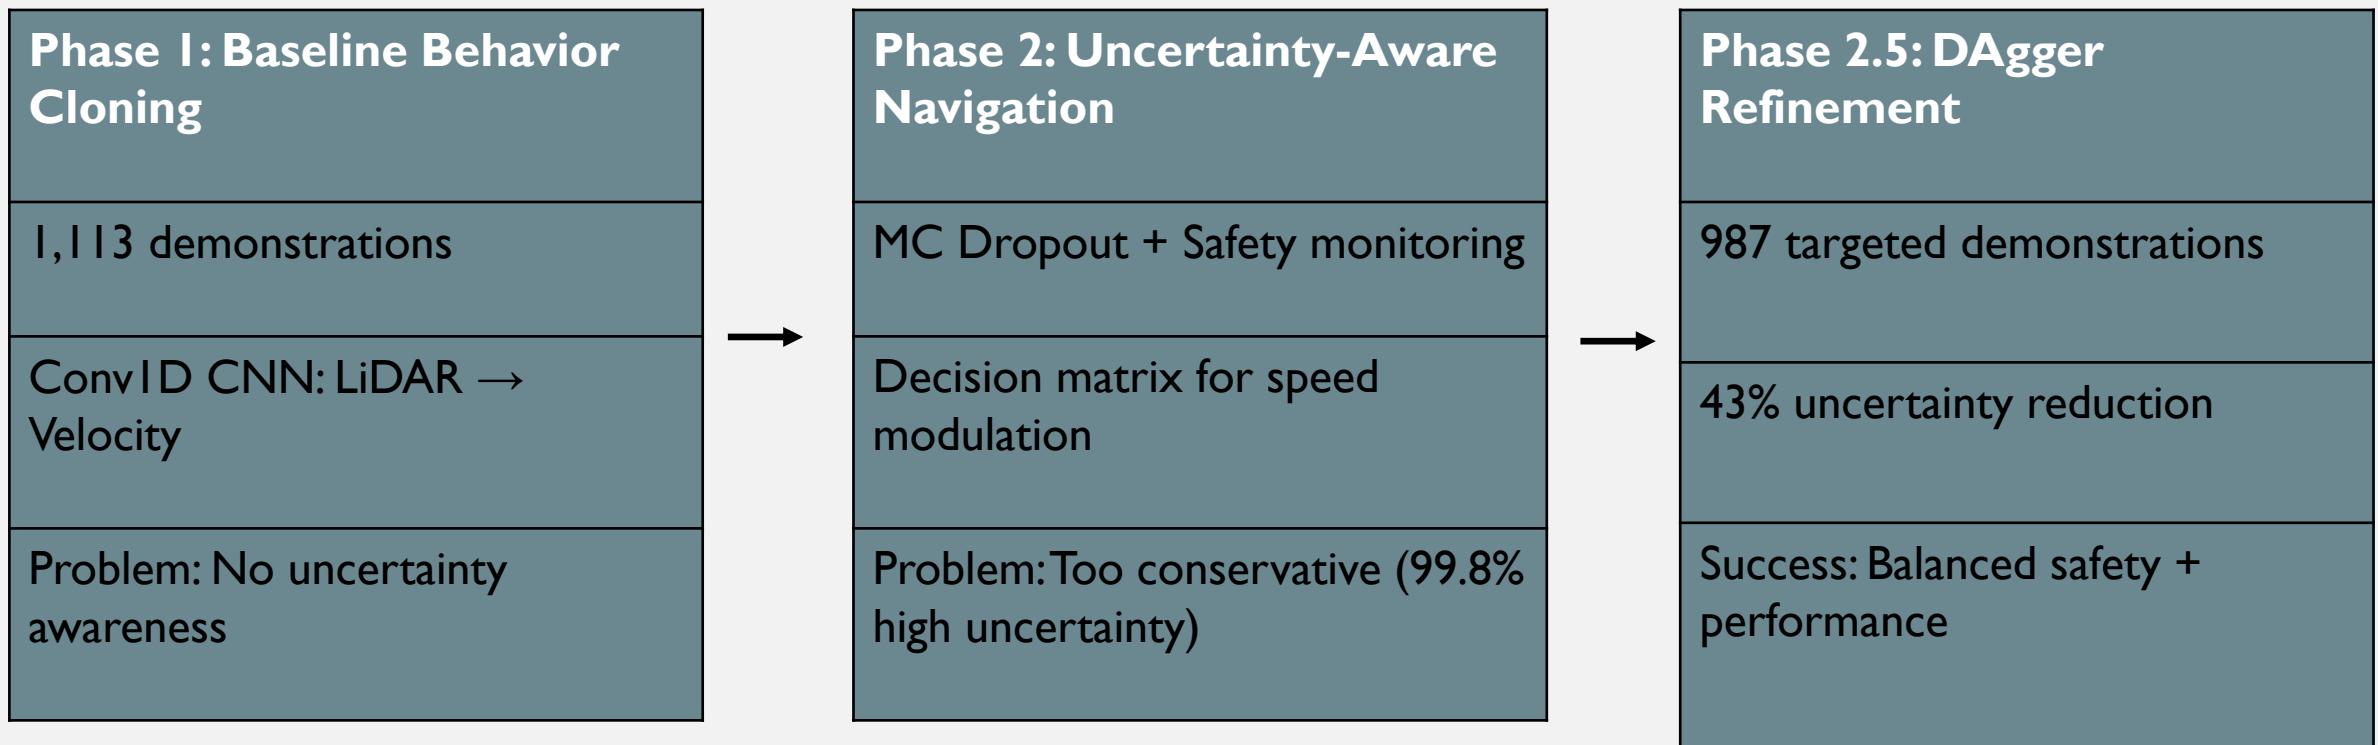
2. Integrate uncertainty with safety monitoring

→ Combine confidence + environmental risk

4. Validate across diverse environments

→ Test generalization without retraining

# THREE-PHASE DEVELOPMENT JOURNEY



# PHASE I: ESTABLISHING THE BASELINE

## Data Collection:

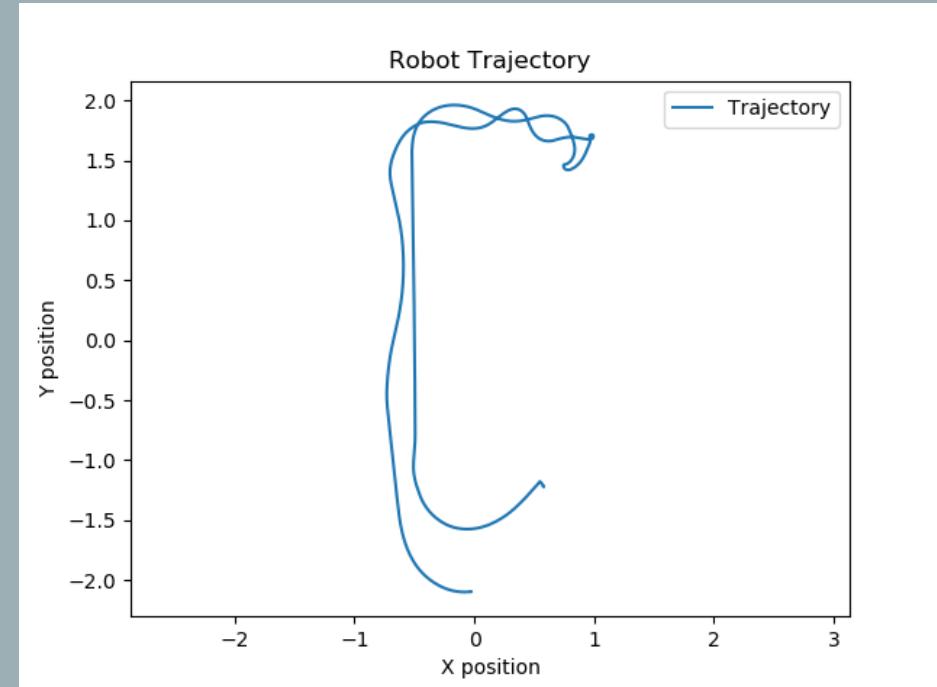
- Manual teleoperation in Gazebo simulation
- TurtleBot3 in turtlebot3\_world environment
- 3 min 17 sec of navigation
- Result: 1,113 synchronized samples

## Network Architecture:

- ConvID CNN: 360-dim LiDAR  $\rightarrow$  2-dim velocity
- 186,178 parameters
- Architecture: Conv  $\rightarrow$  Pool  $\rightarrow$  Conv  $\rightarrow$  Pool  $\rightarrow$  FC  $\rightarrow$  Output

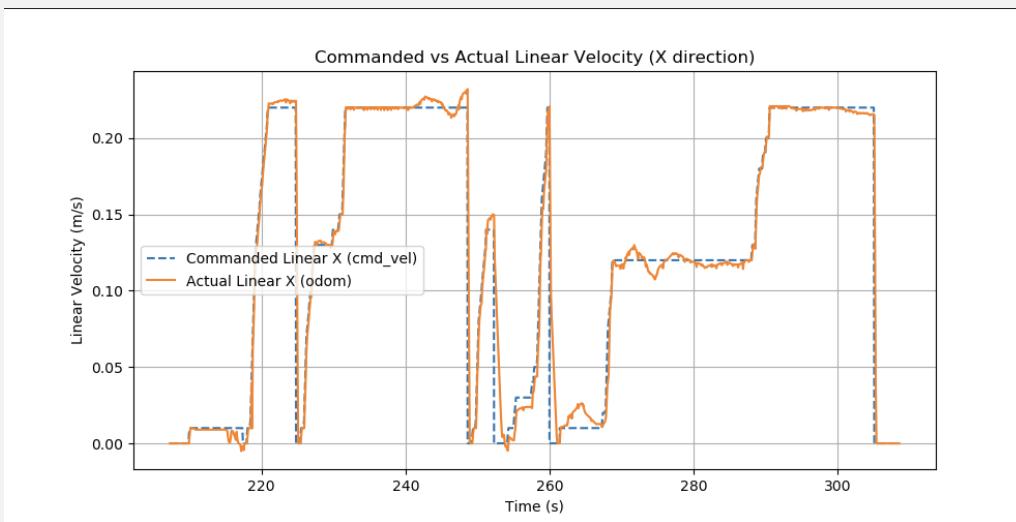
## Training:

- 100 epochs, Adam optimizer
- Test loss: 0.010 (MSE)
- 80/20 train-test split



**Figure I:** Expert demonstration trajectory in 2D space, reconstructed from odometry.

# PHASE I: GOOD PERFORMANCE... BUT LIMITED



**Figure 2:** Phase I: Comparison between predicted velocity commands (red) and ground-truth from expert (blue). The model accurately replicates expert behavior on test data from the same distribution.

## Strengths:

- Accurate tracking of expert behavior
- Fast training (3 minutes)
- Low test loss (0.010)
- Works well in training environment

## Critical Limitation:

- No measure of prediction confidence
- Deterministic outputs regardless of reliability
- Cannot detect unfamiliar situations
- This motivated Phase 2: Add uncertainty awareness

# UNCERTAINTY ESTIMATION: MONTE CARLO DROPOUT

## How MC Dropout Works:

### Traditional Network:

Input → Network → Single Output



**Result:** No Uncertainty X

### MC Dropout (Our Approach):

Input → Network (20×) → Mean

→ Std Dev



**Result:** Uncertainty ✓

**Uncertainty Formula:**  $u = \sqrt{(\sigma^2_{vx} + \sigma^2_{wz})}$

# SAFETY-AWARE CONTROL: THE DECISION MATRIX

**Decision Matrix:** Combining Uncertainty + Environmental Risk

	Low Risk	Med Risk	High Risk
Low Uncertainty	100%	80%	50%
Med Uncertainty	60%	40%	20%
High Uncertainty	30%	20%	STOP

**Uncertainty Levels:**

High:  $u \geq 0.15$

Medium:  $0.10 \leq u < 0.15$

Low:  $u < 0.10$

**Risk Levels:**

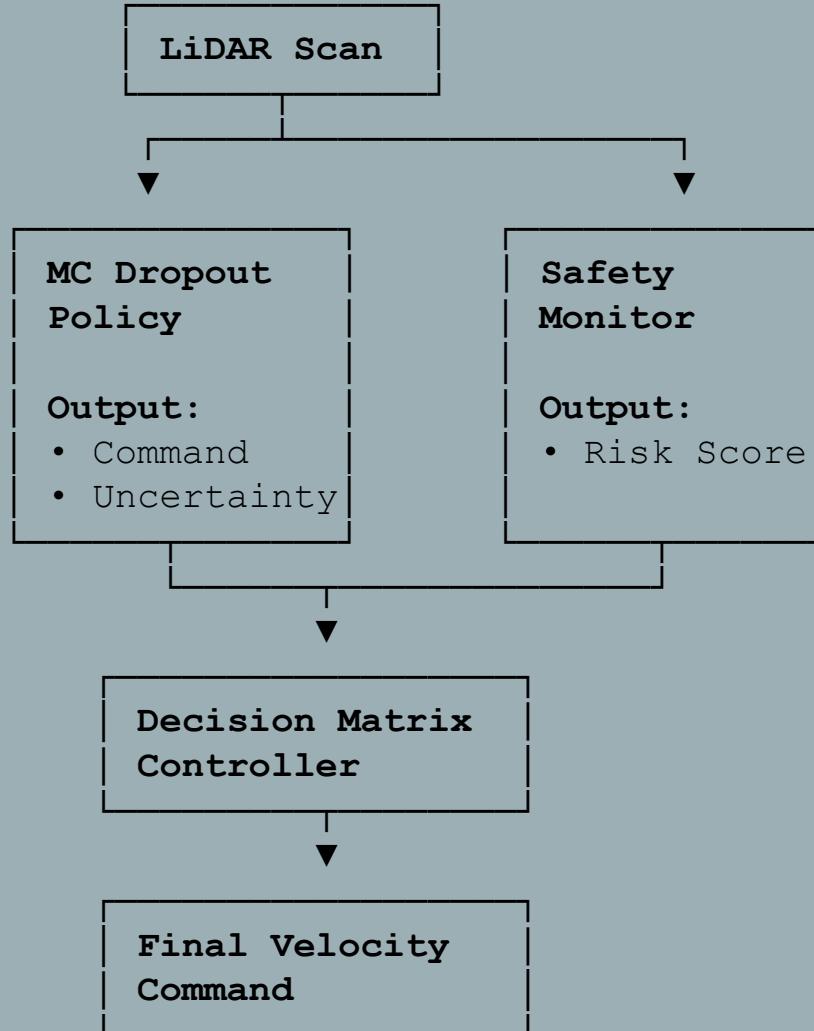
High:  $R \geq 0.85$

Medium:  $0.50 \leq R < 0.85$

Low:  $R < 0.50$

**Key Insight:** Speed adapts to BOTH confidence AND obstacles

# COMPLETE PHASE 2 SYSTEM INTEGRATION



# PHASE 2: SUCCESS... AND A NEW PROBLEM

## Deployment Results:

**SUCCESS:** Uncertainty detection works!

- Mean uncertainty: 0.239
- High uncertainty: 99.8% of the time
- Zero collisions

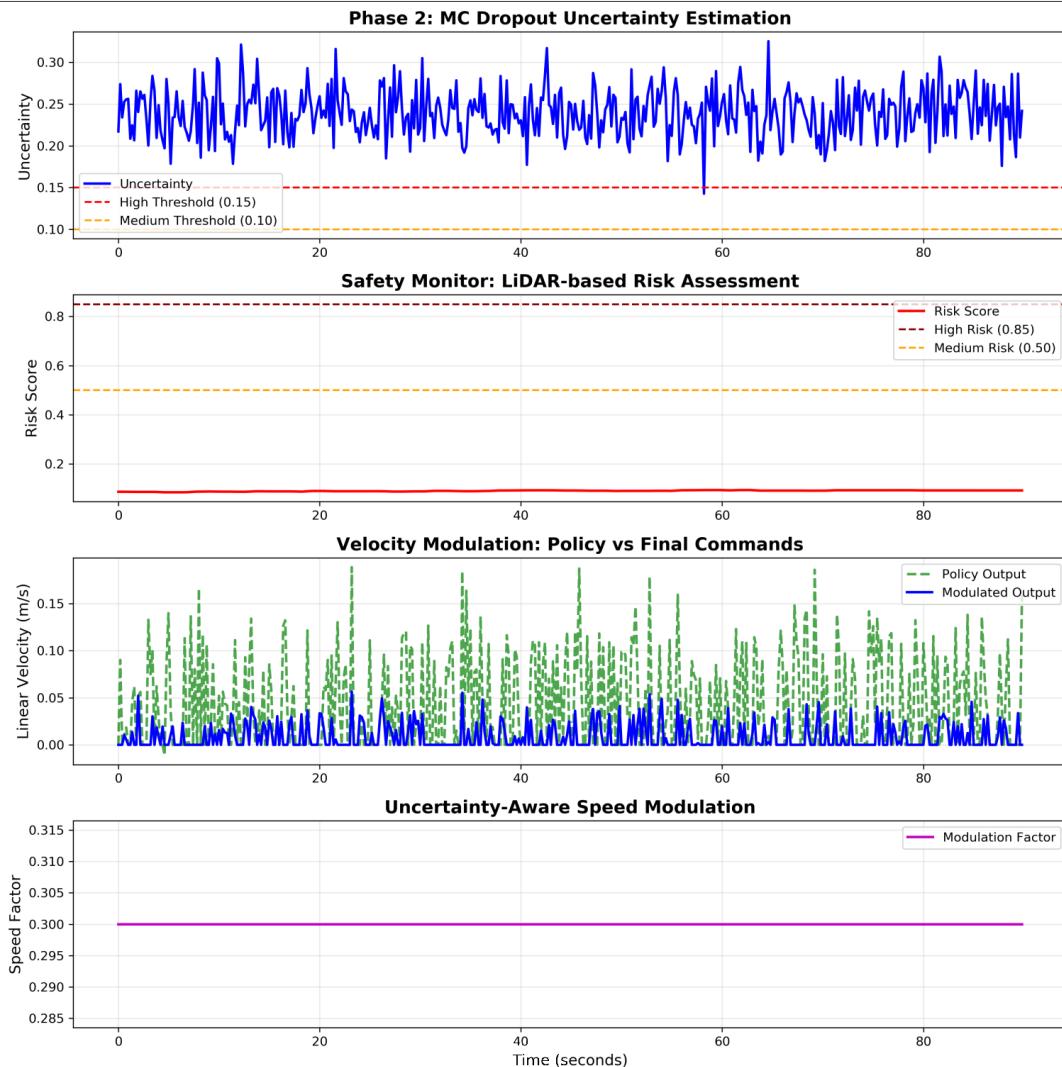
**PROBLEM:** Too conservative

- 70% speed reduction
- Final speed: 0.009 m/s (barely moving)
- Robot is extremely safe but impractical

## Diagnosis:

The robot correctly detected distribution shift—even in the training environment, deployment conditions differ enough to cause uncertainty.

**Solution → Phase 2.5:** Expand training distribution



**Figure 3:** Phase 2 system performance

# TARGETED LEARNING: UNCERTAINTY-GUIDED DAGGER

## Dataset Aggregation (DAgger) Approach:

### Traditional DAgger:

Collect data uniformly →  
Add to dataset →  
Retrain policy →

### Our Uncertainty-Guided DAgger:

Focus on high-uncertainty regions  
Collect targeted demonstrations  
Where  $u > 0.20$

## Data Collection Process:

- Deploy Phase 2 with uncertainty monitoring
- Identify regions where  $u > 0.20$
- Manually demonstrate in those areas
- Collect 987 new samples (3 min 17 sec)
- Merge with Phase 1 data → 2,100 total (+89%)

**Key Insight:** Use uncertainty to guide where to learn

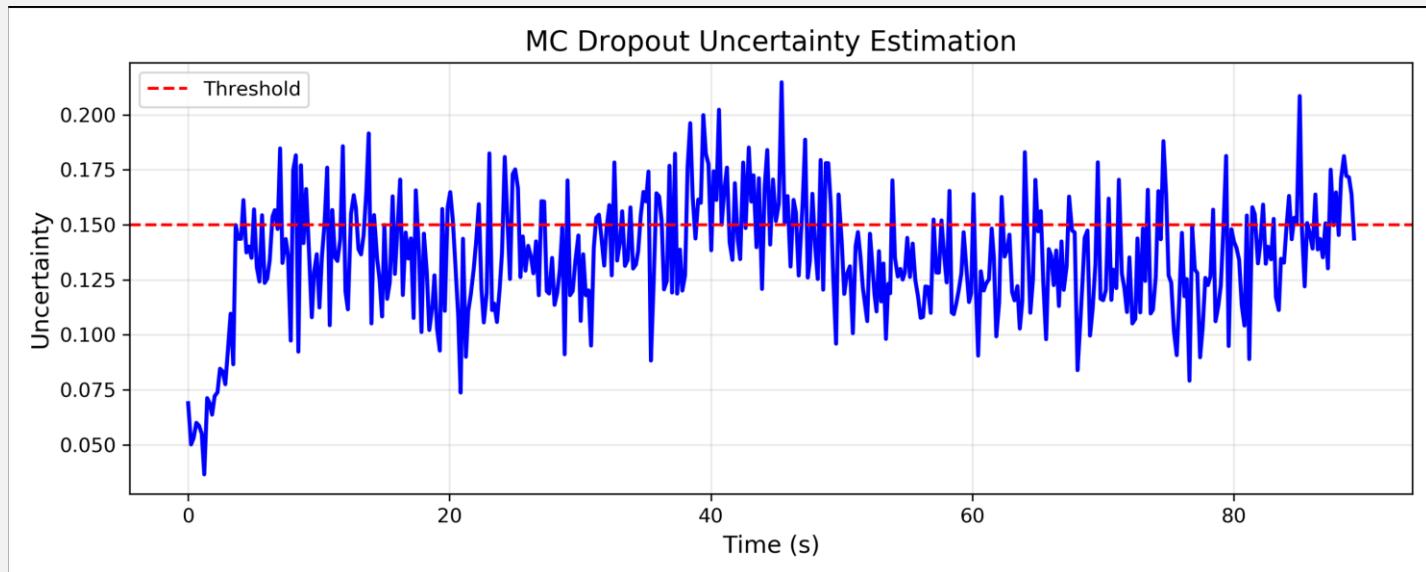
# RETRAINING WITH COMBINED DATASET

## Retraining Details:

- Dataset: 1,113 + 987 = 2,100 samples
- Architecture: Same ConvID CNN (186K parameters)
- Training: 100 epochs, Adam optimizer
- Result: Test loss improved  $0.010 \rightarrow 0.006$  (-40%)

## Training Curves Show:

- ✓ Smooth convergence
- ✓ No overfitting (train/test aligned)
- ✓ Significant improvement in model quality



**Figure 4:** DAgger training curves showing smooth convergence. Final test loss: 0.006, compared to 0.010 in Phase I, indicating improved generalization

# PHASE 2.5: ACHIEVING THE RIGHT BALANCE

## Phase 2.5 Deployment Results:

**SUCCESS:** Balanced behavior achieved!

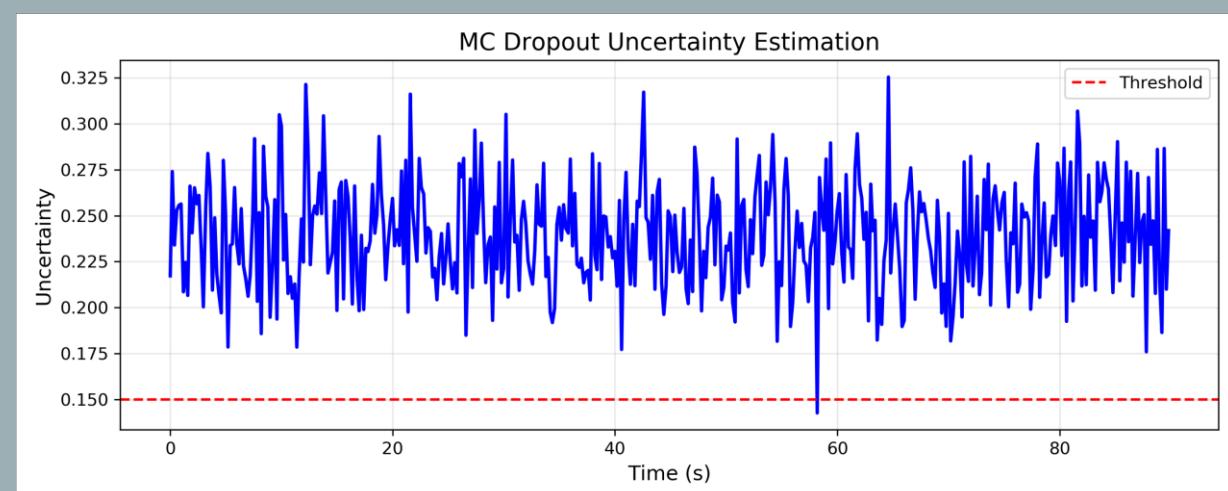
### Uncertainty:

- Mean: 0.135 (down from 0.239) → 43% reduction
- High uncertainty: 28.6% (down from 99.8%)
- Robot is confident 71% of the time

### Performance:

- Final speed: 0.045 m/s (up from 0.009 m/s)
- 5× faster navigation
- 43% speed reduction (down from 70%)
- Still zero collisions

**Result:** Safe when uncertain, efficient when confident



**Figure 5:** Phase 2.5 results after DAgger refinement

## PHASE 2 VS PHASE 2.5: THE NUMBERS

Metric	Phase 2 (Before)	Phase 2.5 (After)	Improvement
Mean Uncertainty	$0.239 \pm 0.029$	$0.135 \pm 0.027$	-43%
% High Uncertainty ( $u > 0.15$ )	99.8%	28.6%	-71%
Speed Reduction	70.0%	43.4%	-38% less conservative
Policy Output Speed	0.042 m/s	0.089 m/s	+112%
Final Navigation Speed	0.009 m/s	0.045 m/s	+400% (5x)
Min Uncertainty Observed	0.143	0.037	-74%
Test Loss (MSE)	0.010	0.006	-40%

- Uncertainty reduced by 43%, from 0.239 to 0.135
- High-uncertainty events dropped from 99.8% to 28.6%, a 71% reduction
- Navigation speed increased 5x, from 0.009 to 0.045 m/s
- Behavior is now adaptive: modulation factor varies dynamically rather than staying constant at 30%

The system achieved the desired balance: cautious when genuinely uncertain (28.6% of the time), but confident and efficient the rest of the time.

# WELL-CALIBRATED UNCERTAINTY: THE ROBOT "KNOWS"

## Uncertainty Distribution Analysis:

### Before DAgger:

Always above threshold

Mean: 0.24

Min: 0.14

High 99.8% of time

### After DAgger:

Frequently below threshold

Mean: 0.14

Min: 0.04 (very confident!)

High 28.6% of time

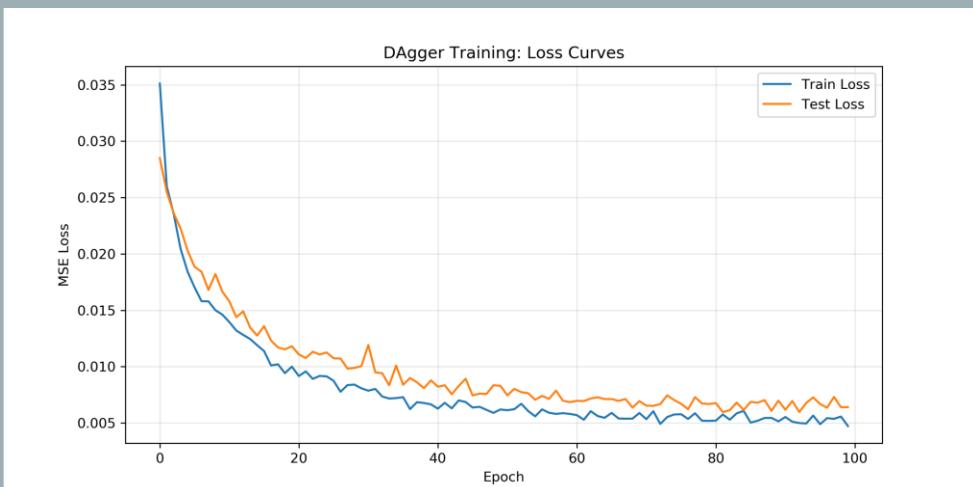


Figure 6: DAgger Training Loss Curves

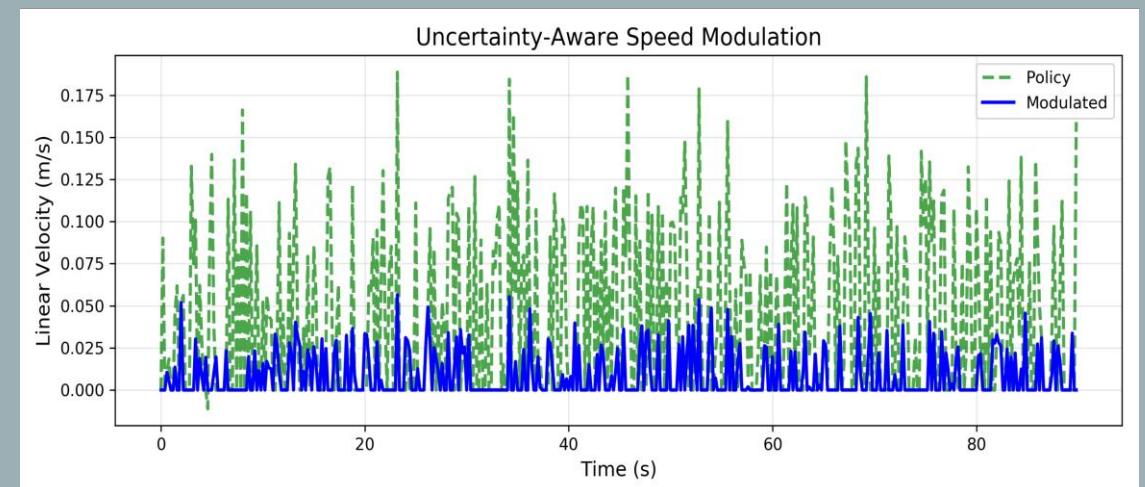


Figure 7: Phase 2.5 (After DAgger)

# DOES IT GENERALIZE? TESTING IN 3 ENVIRONMENTS

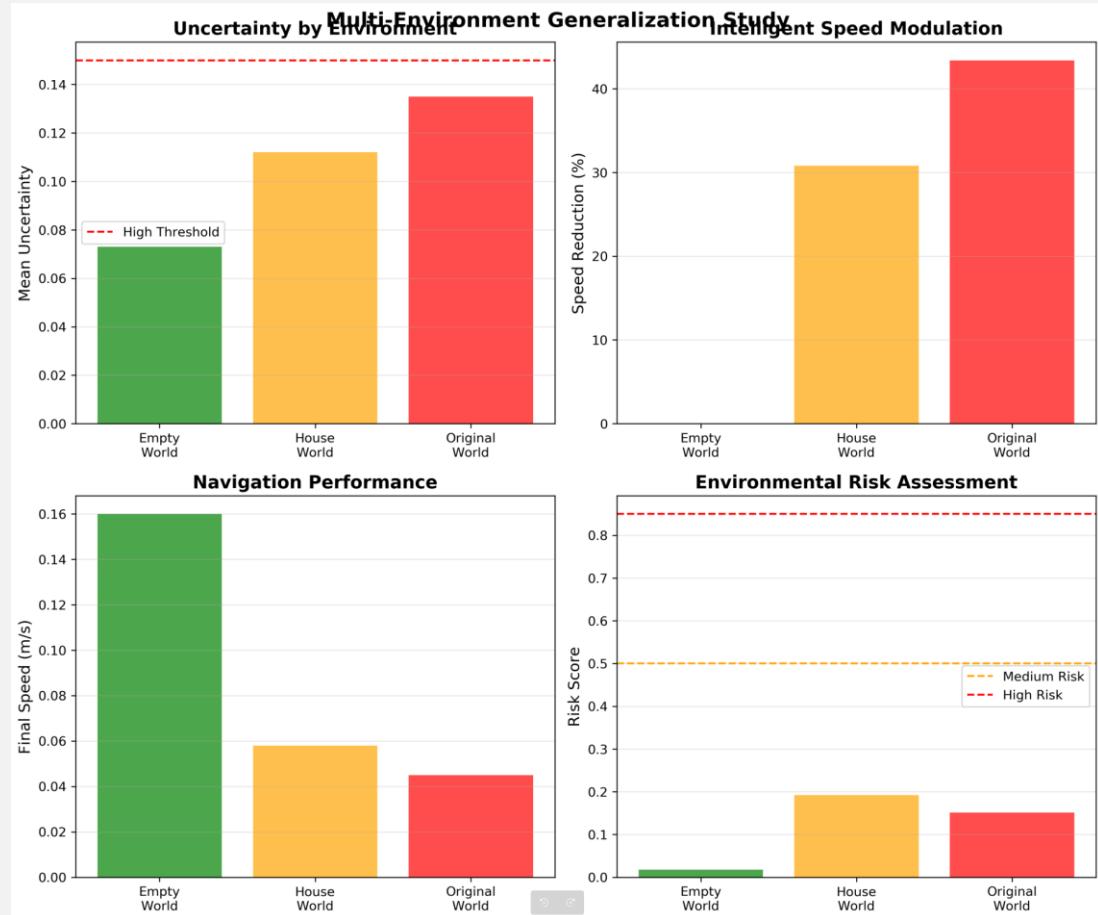
## Test Environments:

- Empty World → Simplest (open space, no obstacles)
- House World → Moderate (furniture, multiple rooms)
- Original World → Most complex (training environment)

## Results:

Environment	Uncertainty	Speed	Reduction
Empty World	0.073	0.160 m/s	0%
House World	0.112	0.058 m/s	31%
Original World	0.135	0.045 m/s	43%

**Key Finding:** Uncertainty correlates with complexity!  
Robot adapts intelligently without retraining

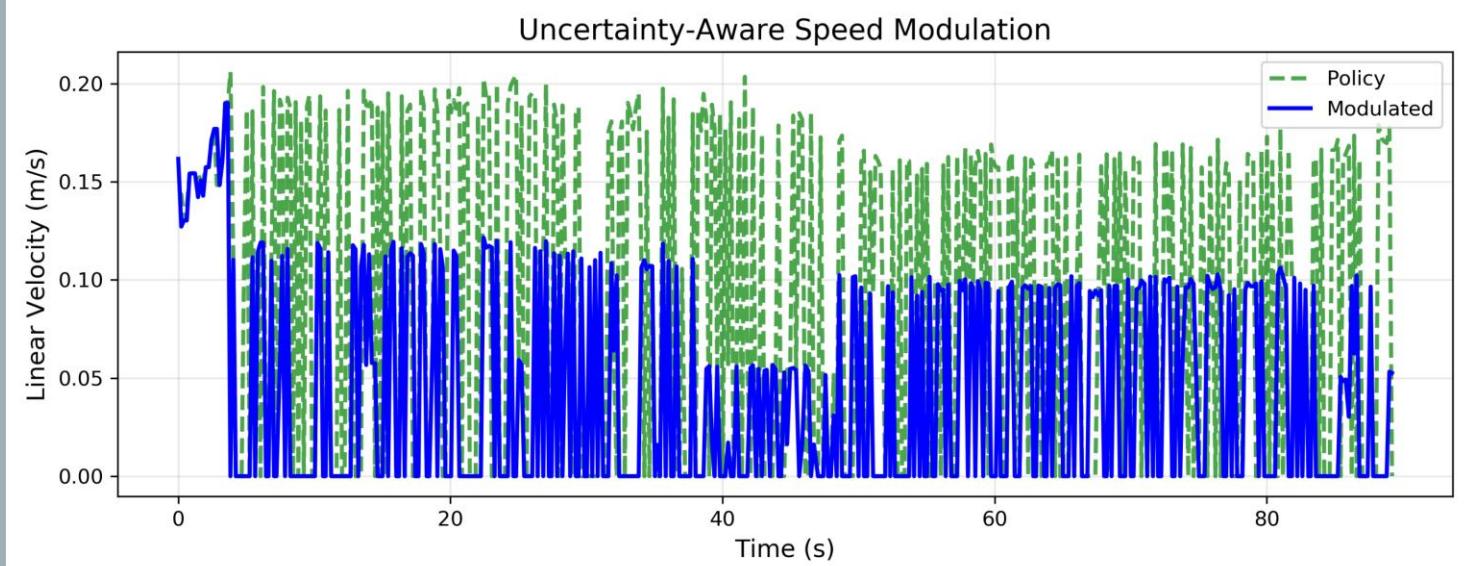


**Figure 8:** Multi-environment performance

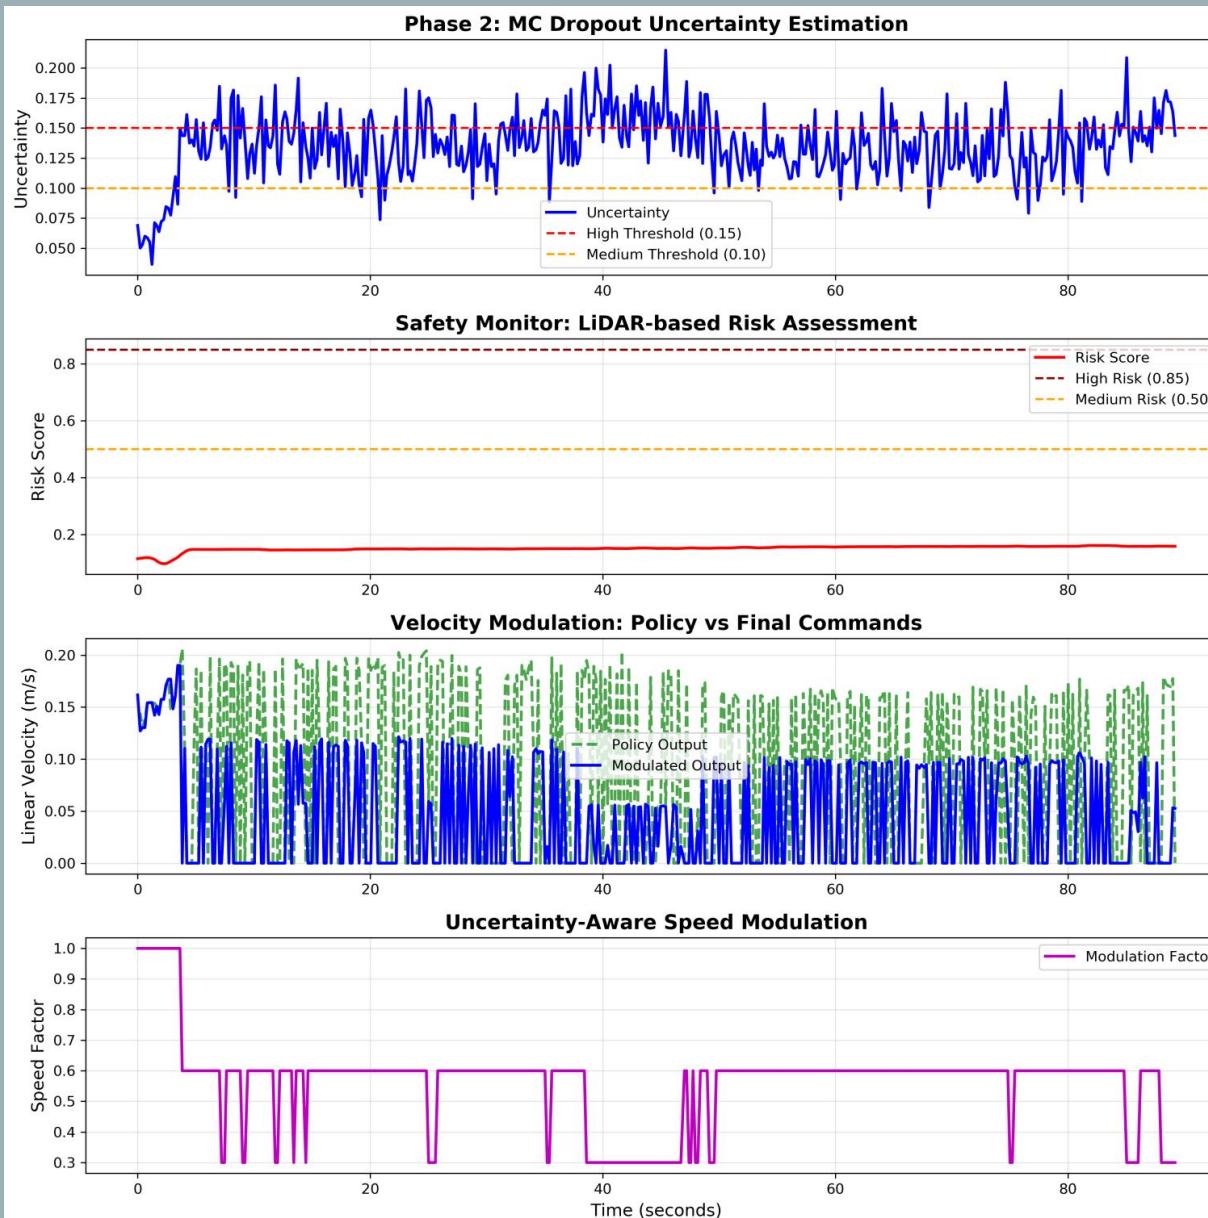
# SPEED MODULATION: BEFORE VS AFTER

## Before DAgger (Phase 2):

- Policy wants: 0.05-0.18 m/s (green dashed)
- Actually executes: 0.01-0.05 m/s (blue solid)
- Huge gap → heavy modulation
- Robot wants to move but can't



**Figure 10:** Phase 2 (Before DAgger)



## After DAgger (Phase 2.5):

- Policy wants: 0.05-0.20 m/s
- Actually executes: 0.05-0.12 m/s
- Much closer tracking
- Robot moves efficiently when confident

## Interpretation:

Lower uncertainty → less aggressive modulation  
→ practical navigation speeds

**Figure II: Phase 2.5 (After DAgger)**

# COMPONENT CONTRIBUTION ANALYSIS

Table 6: Ablation Study: Component Contributions

Configuration	Speed (m/s)	Uncertainty	Reduction	Description
1. Baseline BC	0.042	N/A	0%	Phase 1 policy, no uncertainty
2. BC + Uncertainty	0.045	0.135	15%	MC Dropout only
3. BC + Safety	0.080	N/A	10%	LiDAR safety only
4. Full - DAgger (P2)	0.009	0.239	70%	All components, high $u$
5. Full + DAgger (P2.5)	0.045	0.135	43.4%	Complete system

## Key Findings:

- ✓ Each component contributes measurably
- ✓ DAgger is ESSENTIAL for balance
- ✓ Integration matters (uncertainty + safety together)

# REAL-TIME PERFORMANCE & PRACTICAL VIABILITY

## Performance Metrics:

## Real-Time Operation:

- ✓ Update rate: 5 Hz (200ms per cycle)
- ✓ MC Dropout: 20 forward passes
- ✓ CPU only (no GPU required)
- ✓ Intel i7, 16GB RAM

## Safety Record:

- ✓ Total test time: 15+ minutes across all experiments
- ✓ Collision events: 0 (perfect record)
- ✓ Min obstacle distance: Always  $> 0.05m$

## Data Efficiency:

- ✓ Phase I collection: 3:17 minutes
- ✓ DAgger collection: 3:17 minutes
- ✓ Total manual effort: < 7 minutes
- ✓ Training time: 3 minutes per phase

- **Deployment Ready:** ✓ Practical for resource-constrained robots

# MAIN CONTRIBUTIONS OF THIS WORK

## 1. Complete Uncertainty-Aware Navigation System

- Real-time MC Dropout (5 Hz)
- Integrated with LiDAR safety monitoring
- Practical deployment on standard hardware

## 4. Comprehensive Validation

- Multi-environment testing
- Rigorous ablation study (5 configurations)
- Each component's contribution quantified

## 2. Distribution Shift Detection & Resolution

- Uncertainty reveals distribution shift ( $99.8\% \rightarrow 28.6\%$ )
- DAgger with uncertainty-guided collection
- 43% uncertainty reduction, 5 $\times$  speed improvement

## 5. Reproducible Methodology

- Complete implementation details
- Data-efficient (7 min collection)
- Framework for future work

## 3. Well-Calibrated Uncertainty Across Environments

- Correlates with complexity ( $0.073 \rightarrow 0.135$ )
- Validates "knowing when you don't know"
- Generalizes without retraining

# LIMITATIONS AND NEXT STEPS

## Current Limitations:

- Simulation only (Gazebo) - needs real robot validation
- Limited baseline comparisons - no ensemble methods tested
- Single platform (TurtleBot3) - needs broader validation
- Three environments - could test in more diverse settings

## Future Work:

- Priority 1: Real Robot Deployment  
Deploy on physical TurtleBot3 to validate sim-to-real transfer
- Priority 2: Baseline Comparisons  
Compare vs ensembles, Bayesian NNs, other uncertainty methods
- Priority 3: Extended Testing  
10+ environments, long-term deployment, dynamic obstacles
- Priority 4: Active Learning  
Robot autonomously identifies uncertain regions and requests help

# FROM DETERMINISTIC TO UNCERTAINTY-AWARE TO BALANCED

Phase 1	Phase 2	Phase 2.5
Deterministic	Overly Conservative	Balanced
No awareness	High uncertainty	Practical
0.042 m/s	0.009 m/s	0.045 m/s
	99.8% uncertain	28.6% uncertain

## The Value of the Journey:

- ✓ Phase 2's "failure" was actually success—it revealed the problem
- ✓ Systematic problem identification → targeted solution → validation
- ✓ Each phase taught us something essential

# Thank You!

Shubham Yogesh Jangle

Master's Student, Robotics

University of California, Riverside

[sjang041@ucr.edu](mailto:sjang041@ucr.edu)

## Key Results:

- ✓ 43% uncertainty reduction through targeted learning
- ✓ 5× navigation speed improvement
- ✓ Well-calibrated across 3 environments
- ✓ Zero collisions, real-time operation (5 Hz)